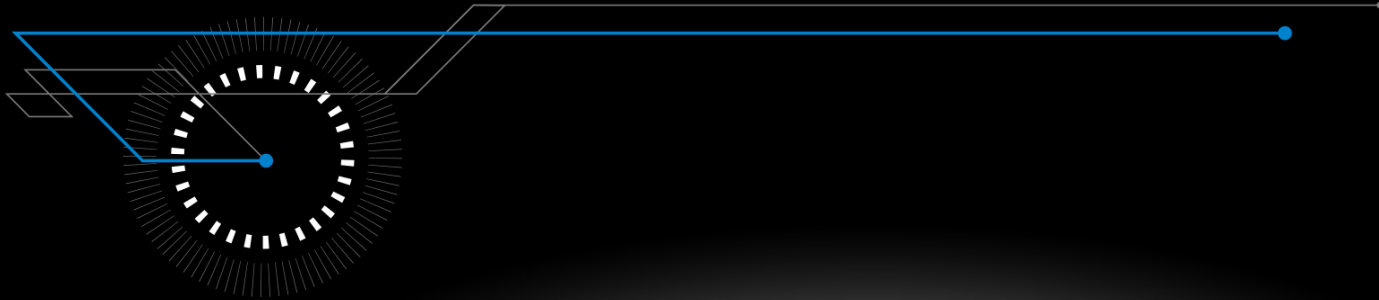


Scalable Microservices at Netflix. Challenges and Tools of the Trade

AGENDA

- Netflix – background and evolution
 - Monolithic Apps
 - Characteristics
 - What are Microservices?
 - Microservices
 - Why?
 - Challenges
 - Best practices
 - Tools of the trade
 - InterProcess Communication
 - Takeaways
-

Netflix - Evolution



Netflix - Evolution

- Old DataCenter (2008)
- Everything in one WebApp (.war)
- AWS Cloud (~2010)
- 100s of Fine Grained Services

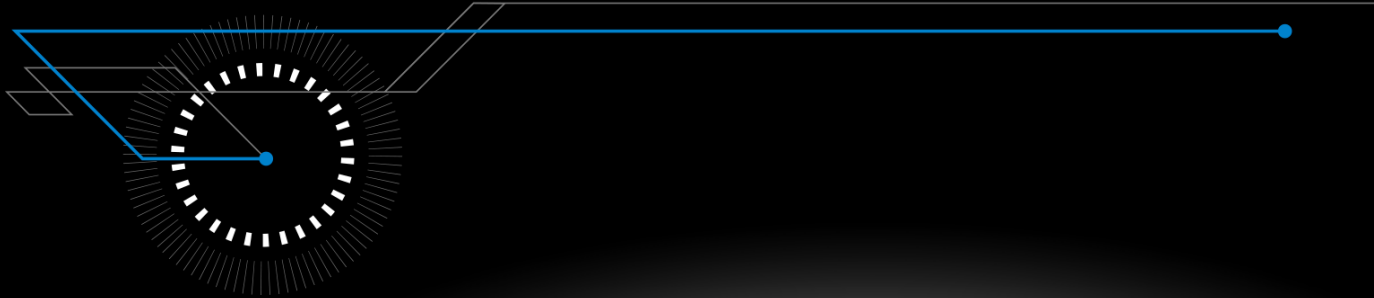



Netflix Scale

- ~ 1/3 of the peak Internet traffic a day
- ~50M subscribers
- ~2 Billion Edge API Requests/Day
- >500 MicroServices
- ~30 Engineering Teams (owning many microservices)



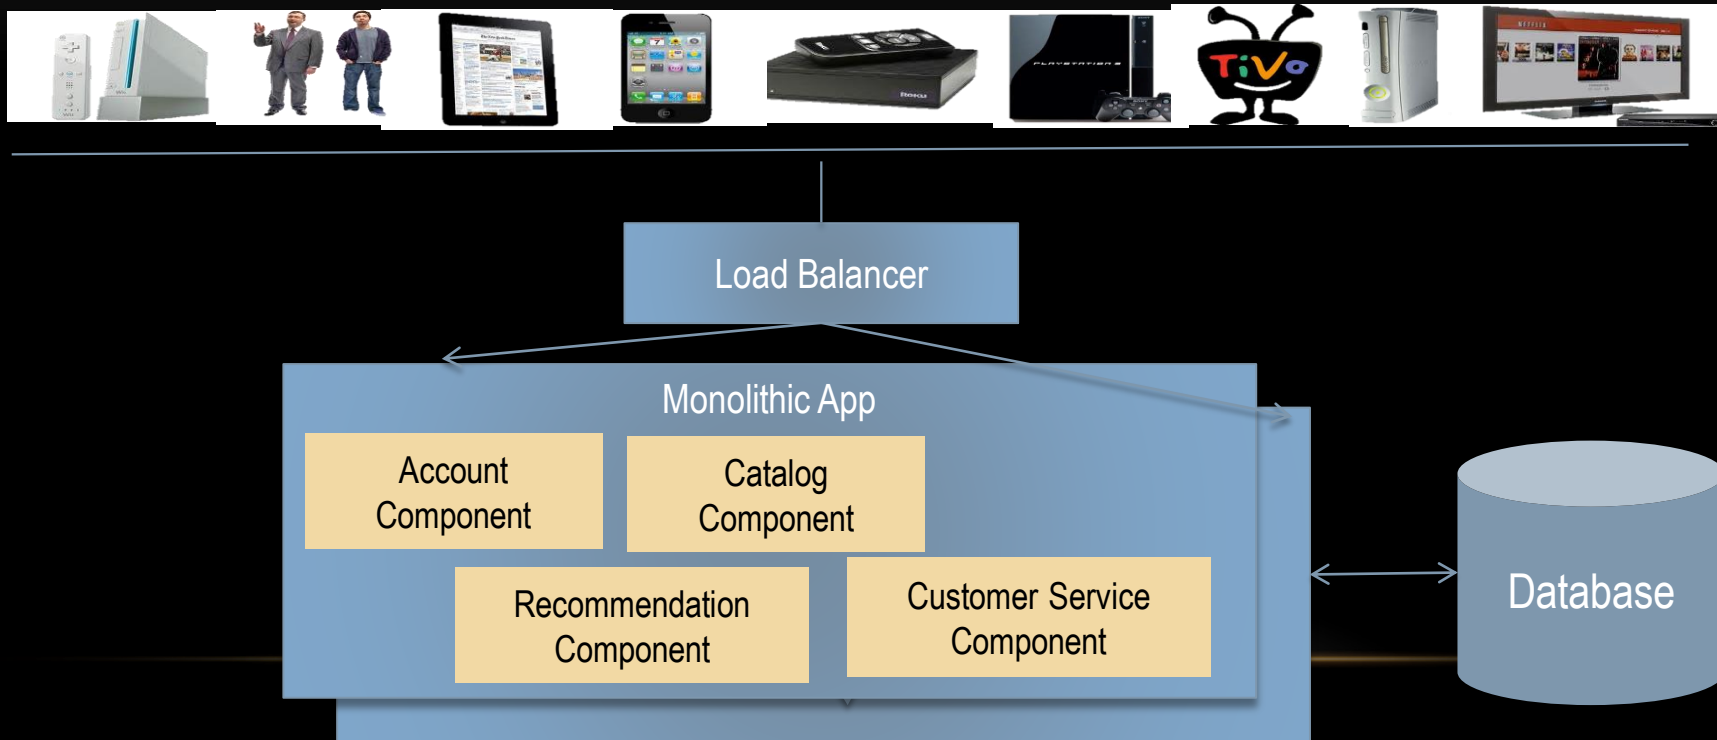
Monolithic Apps





MONOLITHIC APP

Monolithic Architecture



Characteristics

- Large Codebase
- Many Components, no clear ownership
- Long deployment cycles

Pros

- Single codebase
 - Easy to develop/debug/deploy
 - Good IDE support
- Easy to scale horizontally (but can only scale in an “un-differentiated” manner)
- A Central Ops team can efficiently handle

Monolithic App – Evolution

- As codebase increases ...
 - Tends to increase “tight coupling” between components
 - Just like the cars of a train
 - All components have to be coded in the same language






User Accounts

Customer Service

Shopping Cart

Product Catalog



Evolution of a Monolithic App

Monolithic App - Scaling

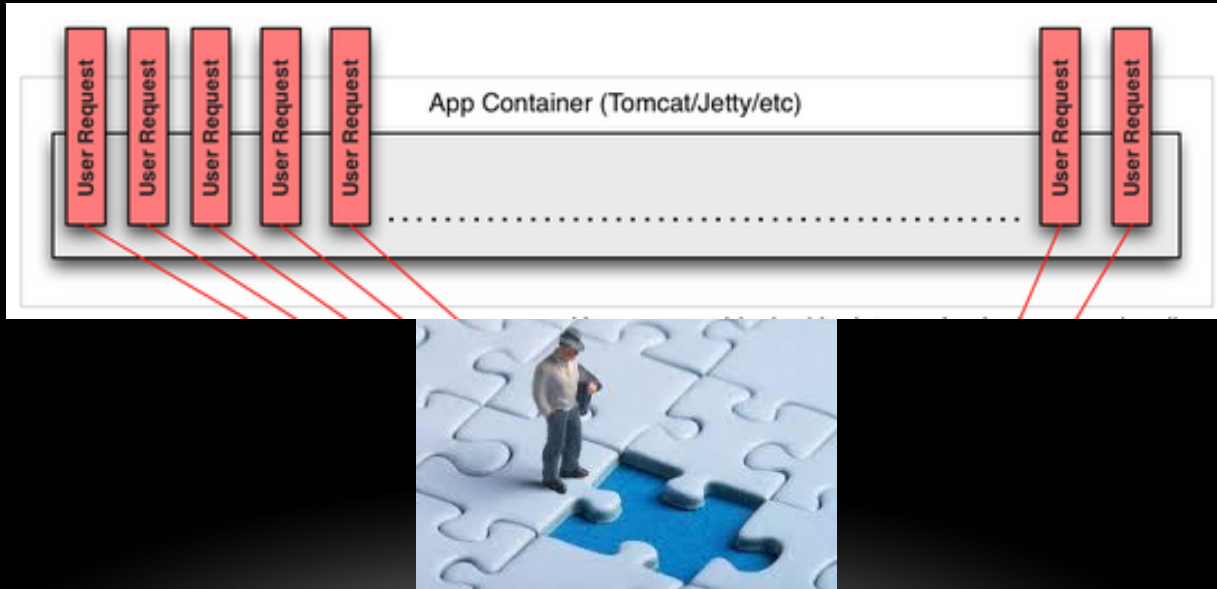
- Scaling is “undifferentiated”
 - Cant scale “Product Catalog” differently from “Customer Service”



AVAILABILITY

Availability

- A single missing “;” brought down the Netflix website for many hours (~2008)





MONOLITHIC APPS – FAILURE & AVAILABILITY



MicroServices


You Think??

TIPPING POINT



Organizational Growth

&



Diverse Functionality

&




Bottleneck in
Monolithic stack

What are MicroServices?



NOT ABOUT ...

- Team size
- Lines of code
- Number of API/EndPoints



CHARACTERISTICS

- Many smaller (fine grained), clearly scoped services
 - Single Responsibility Principle
 - Domain Driven Development
 - Bounded Context
 - Independently Managed
- Clear ownership for each service
 - Typically need/adopt the “DevOps” model



Composability– unix philosophy

- Write programs that do one thing and do it well.
- Write programs to work together

```
tr 'A-Z' 'a-z' < doc.txt | tr -cs 'a-z' '\n' | sort | uniq | comm -23 - /usr/share/dict/words
```



Program to print misspelt words in doc.txt

Comparing Monolithic to MicroServices

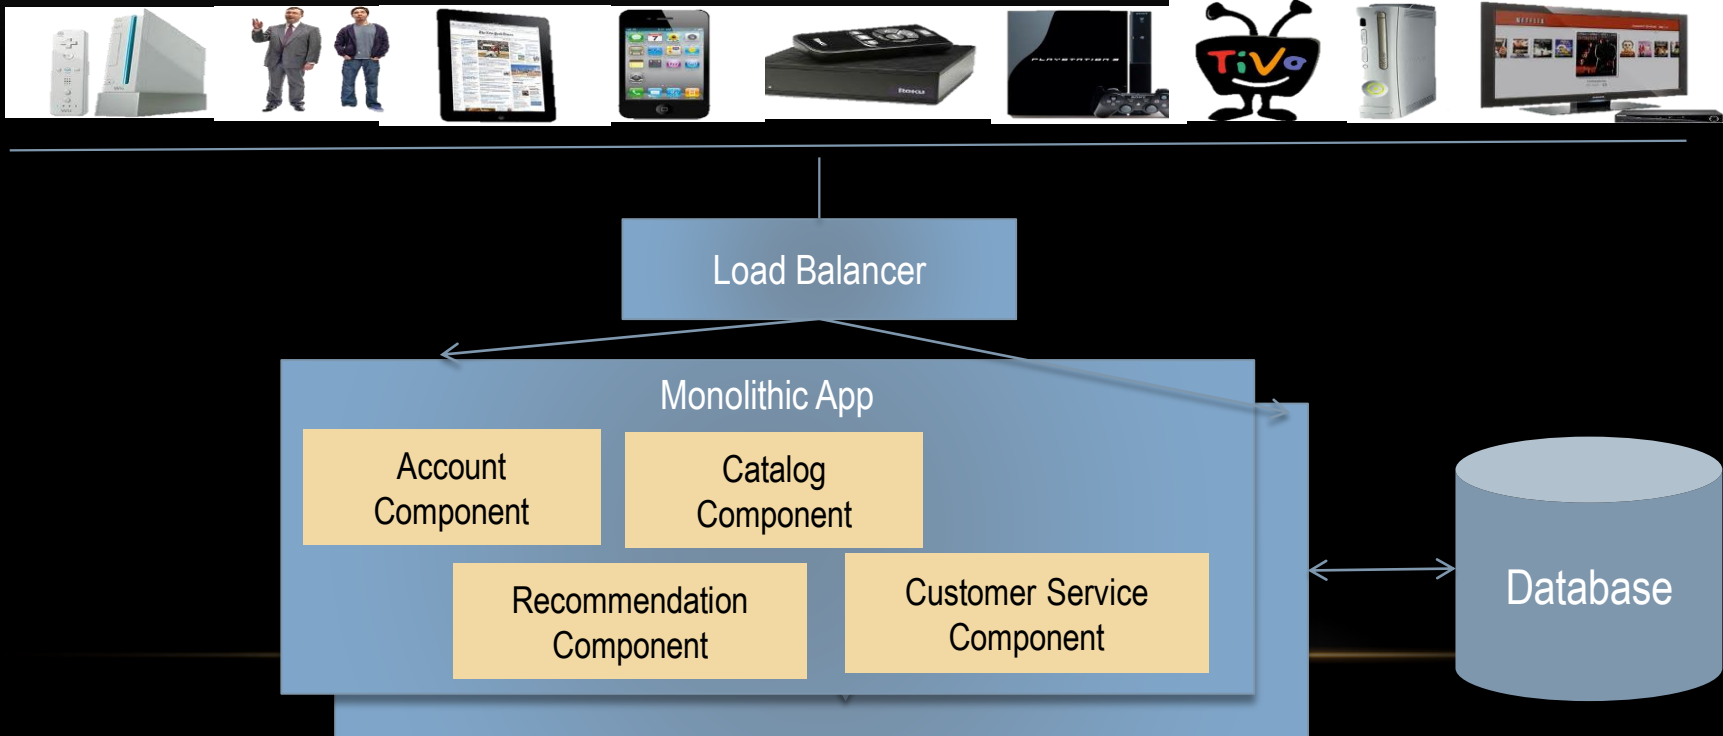


MONOLITHIC APP (VARIOUS COMPONENTS LINKED TOGETHER)

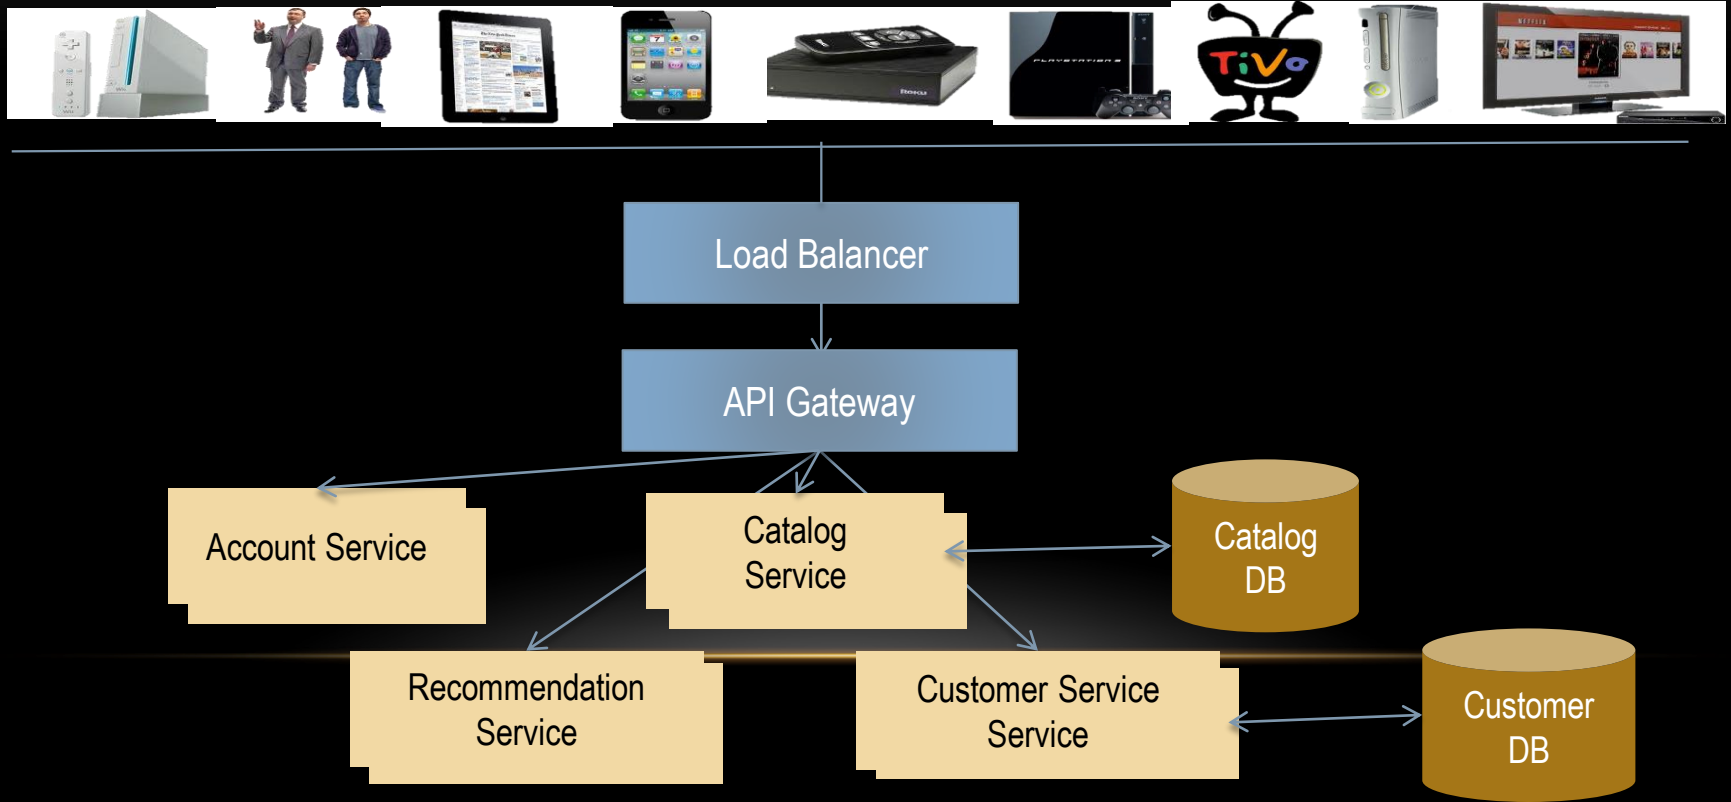


MICROSERVICES – SEPARATE SINGLE PURPOSE SERVICES


Monolithic Architecture (Revisiting)



Microservices Architecture

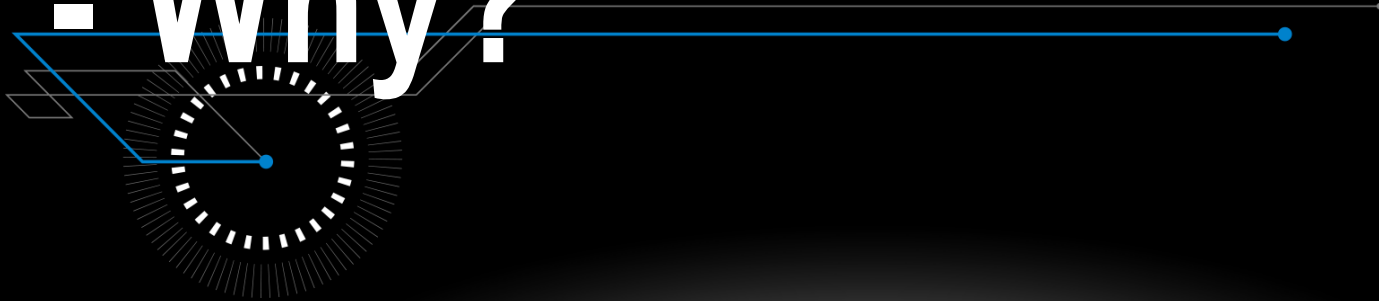


Concept -> Service Dependency Graph



MicroServices

- Why?



WHY?


- Faster and simpler deployments and rollbacks
 - Independent Speed of Delivery (by different teams)
- Right framework/tool/language for each domain
 - Recommendation component using Python?, Catalog Service in Java ..
- Greater Resiliency
 - Fault Isolation
- Better Availability
 - If architected right 😊

MicroServices

- Challenges



CHALLENGES



Can lead to chaos if not designed right ...

OVERALL COMPLEXITY

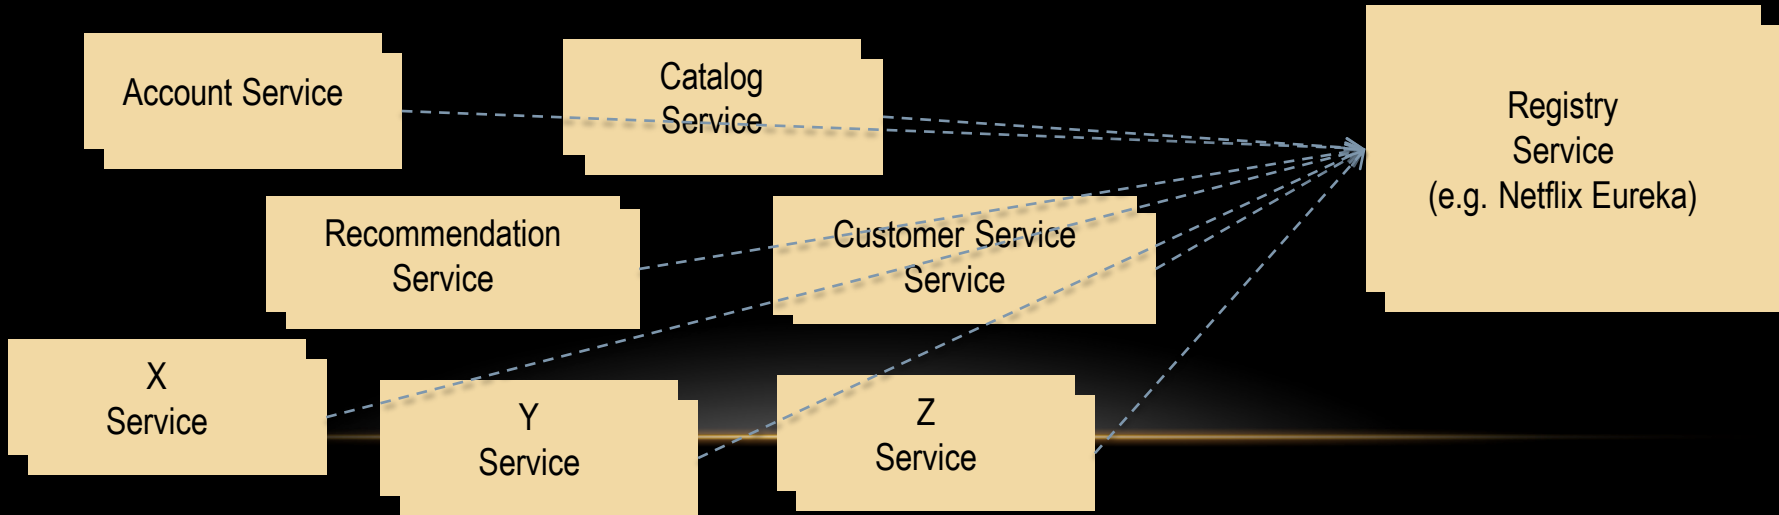
- Distributed Systems are inherently Complex
 - N/W Latency, Fault Tolerance, Retry storms ..
- Operational Overhead
 - TIP: Embrace DevOps Model

SERVICE DISCOVERY

- 100s of MicroServices
 - Need a Service Metadata Registry (Discovery Service)





NETFLIX | OSS



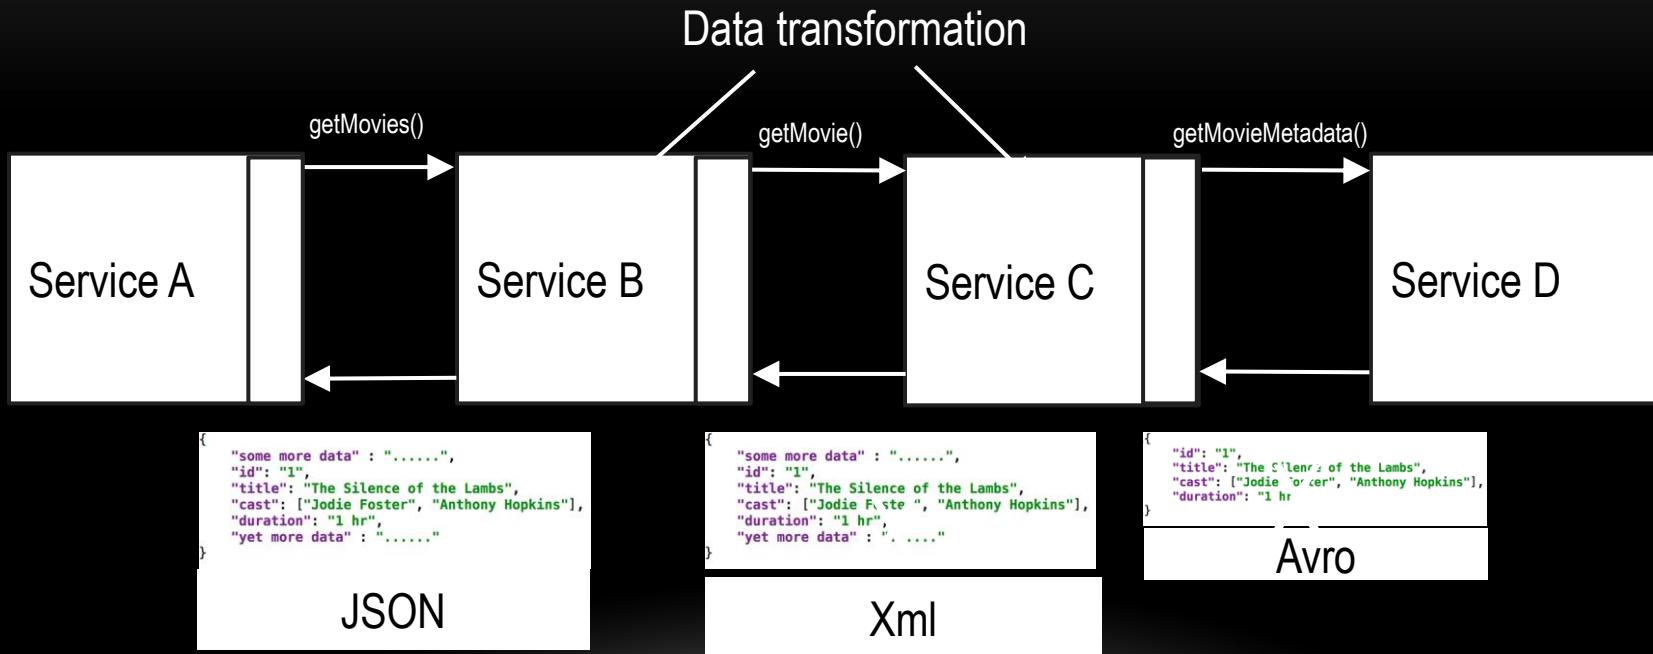
CHATTINESS (AND FAN OUT)

~2 Billion Requests per day on Edge Service

Results in ~20 Billion Fan out requests in ~100 MicroServices




DATA SERIALIZATION OVERHEAD



CHALLENGES - SUMMARY

- Service Discovery
- Operational Overhead (100s of services; DevOps model absolutely required)
- Distributed Systems are inherently Complex
 - N/W Latency, Fault Tolerance, Serialization overhead ..
- Service Interface Versioning, Mismatches?
- Testing (Need the entire ecosystem to test)
- Fan out of Requests -> Increases n/w traffic

Best Practices/Tips



Best Practice -> Isolation/Access

- TIP: In AWS, use Security Groups to isolate/restrict access to your MicroServices

Edit inbound rules [X]

Type ⓘ	Protocol ⓘ	Port Range ⓘ	Source ⓘ	
HTTP ⌵	TCP	80	Anywhere ⌵ 0.0.0.0/0	✕
HTTPS ⌵	TCP	443	Anywhere ⌵ 0.0.0.0/0	✕

[Add Rule] [Cancel] [Save]

Best Practice -> Loadbalancers

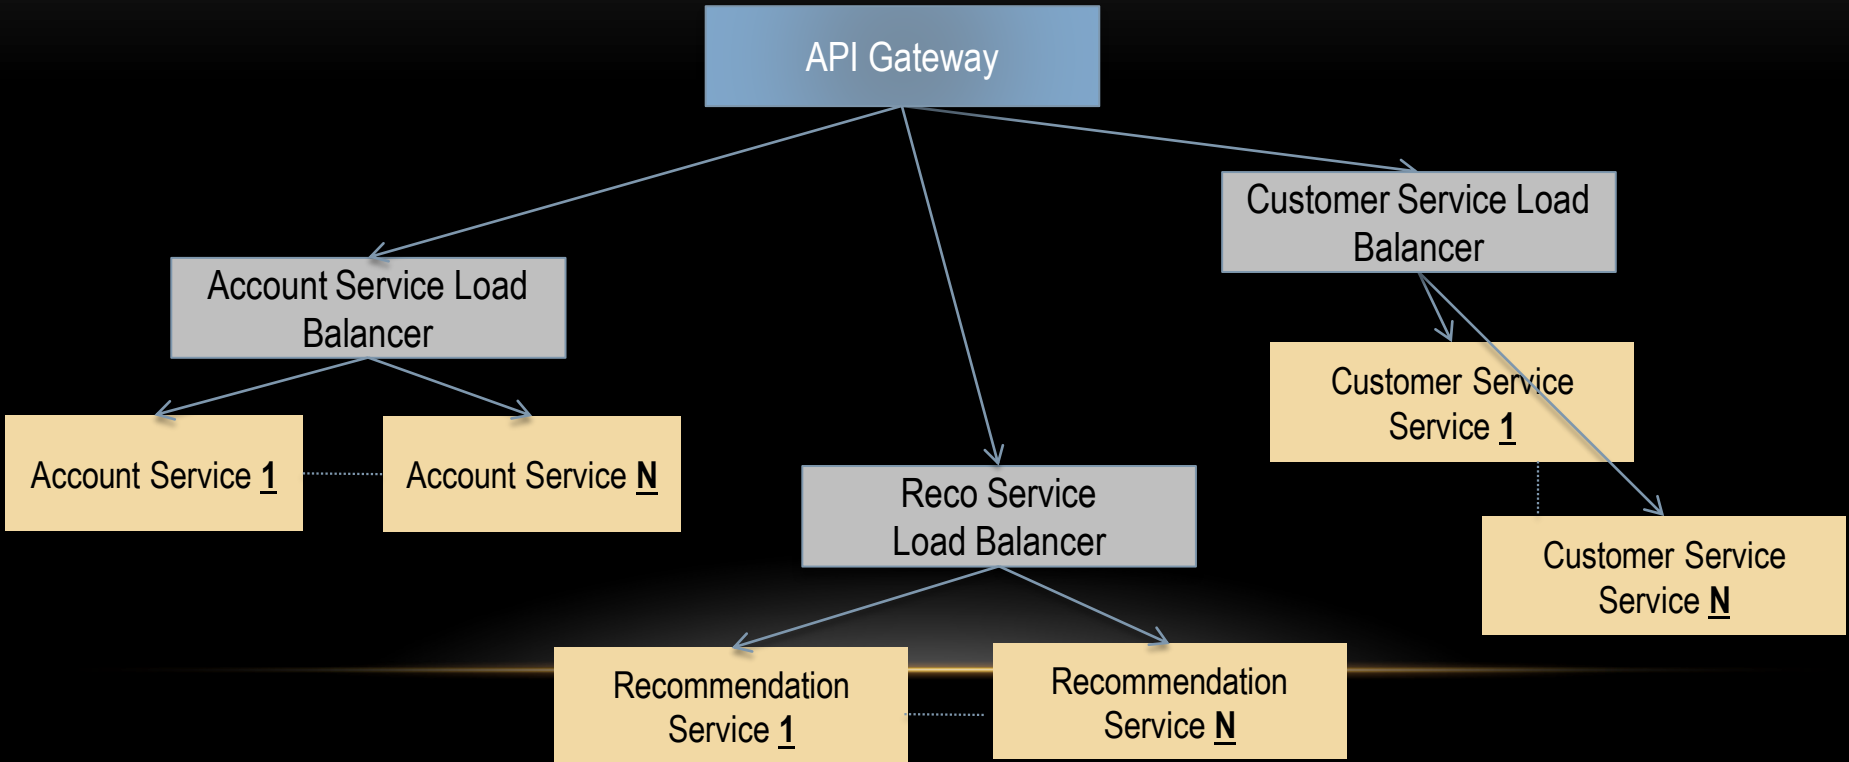
Choice

1. Central Loadbalancer? (H/W or S/W)

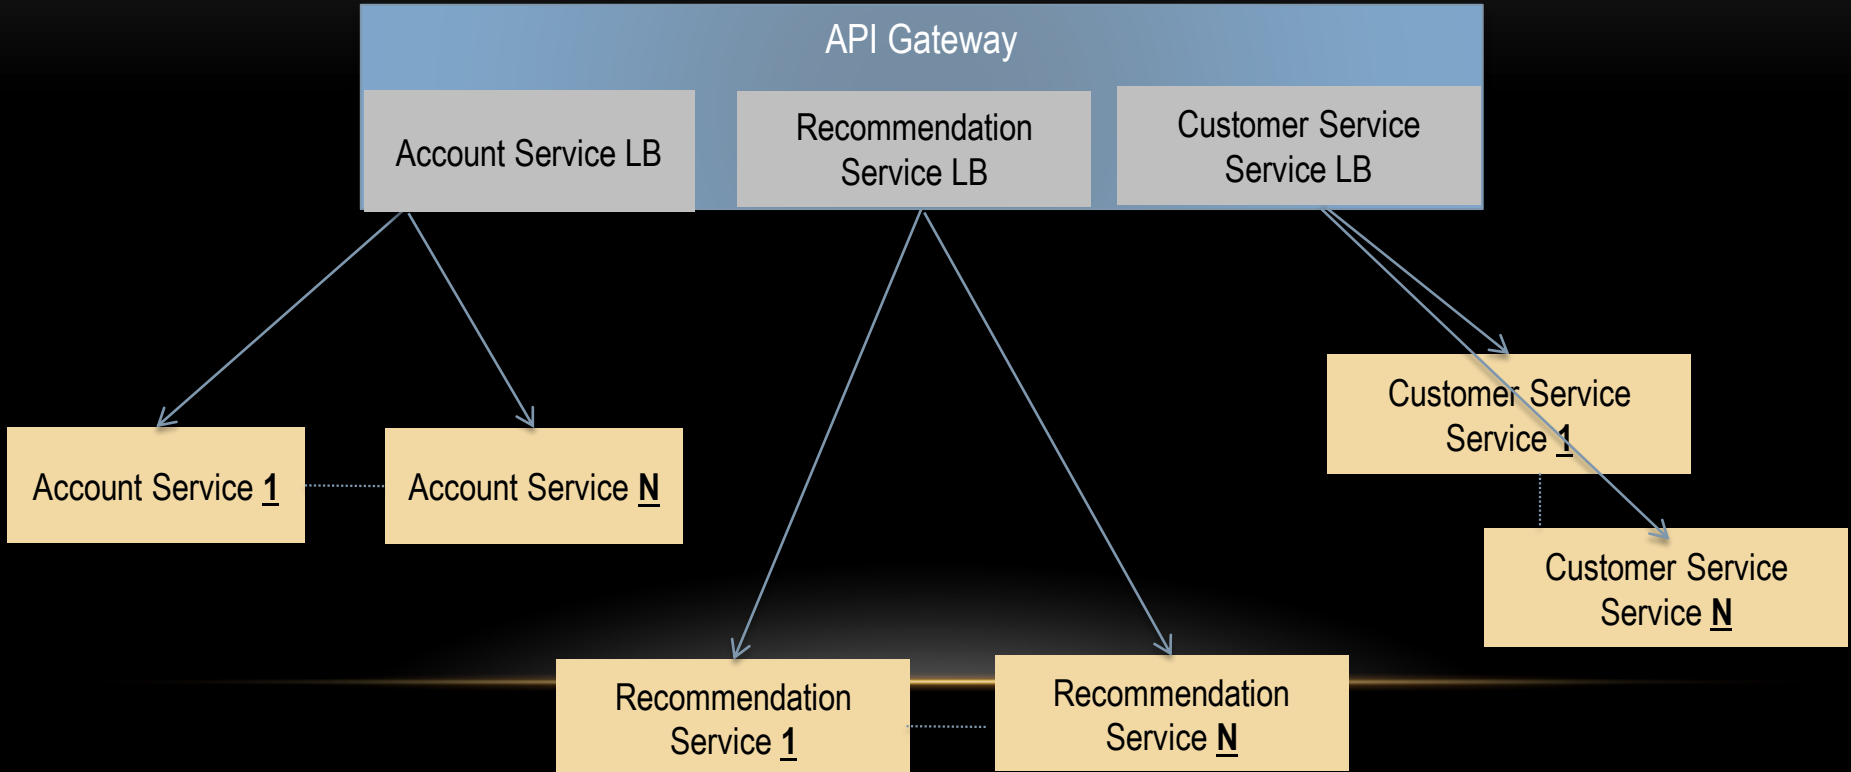
OR

2. Client based S/W Loadbalancer?


Central (Proxy) Loadbalancer



Client Loadbalancer



Client based Smart Loadbalancer



Use Ribbon (<http://github.com/netflix/ribbon>)

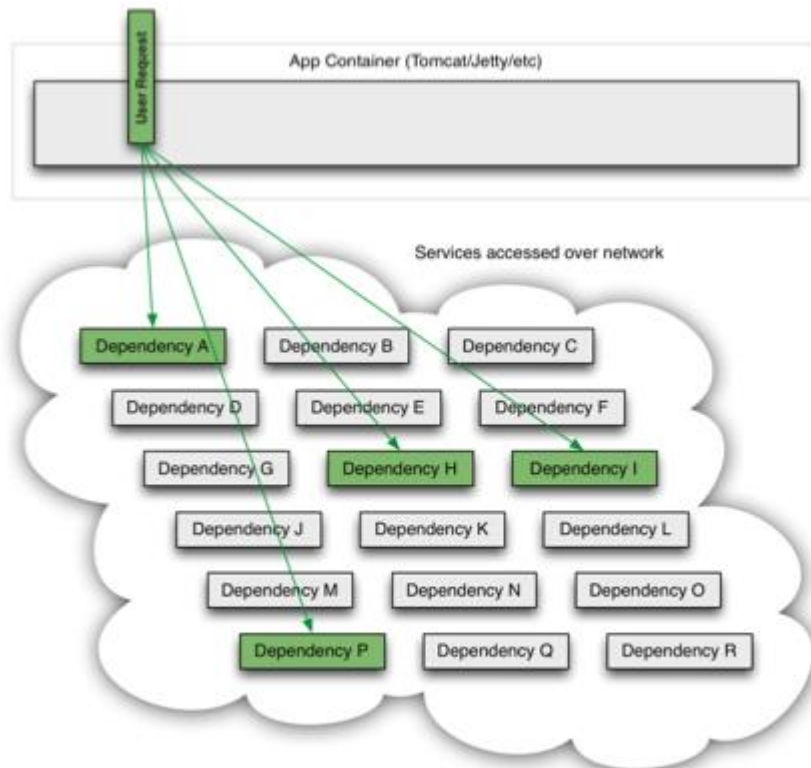
Best Practice -> LoadBalancers

- TIP: Use Client Side Smart LoadBalancers

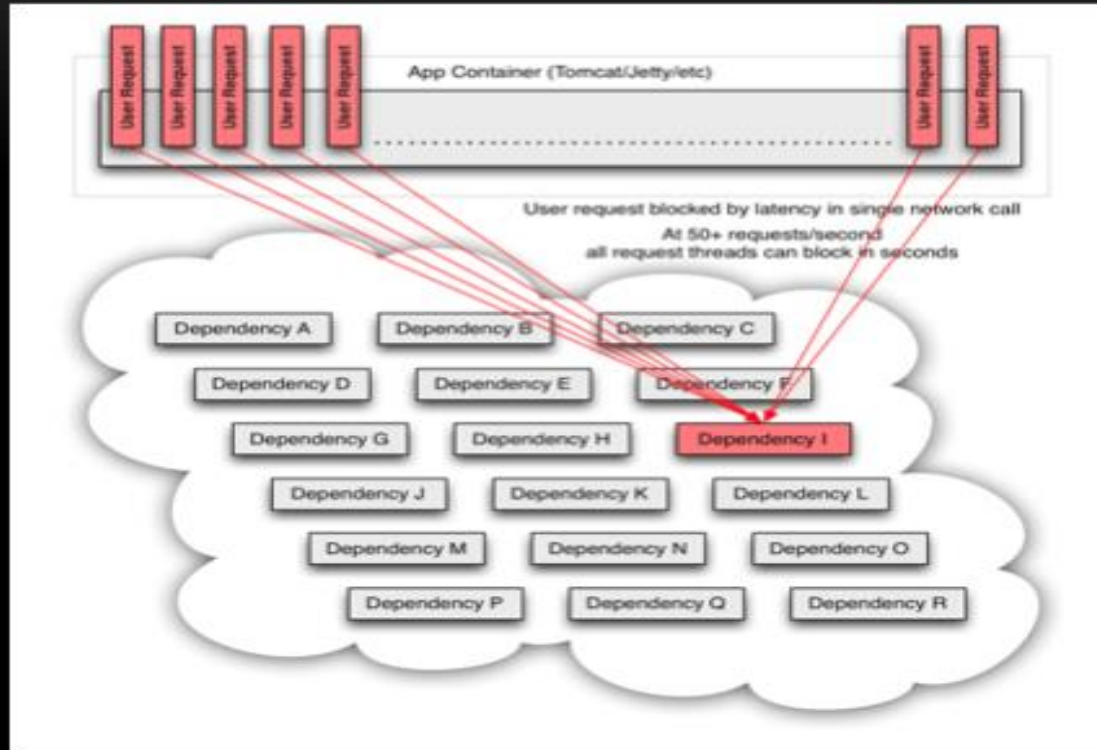
BEST PRACTICES CONTD..

- Dependency Calls
 - Guard your dependency calls
 - Cache your dependency call results
 - Consider Batching your dependency calls
 - Increase throughput via Async/ReactiveX patterns
-

Dependency Resiliency



Service Hosed!!



A single “bad” service can still bring your service down

AVAILABILITY

MicroServices does **not** automatically mean better Availability
- Unless you have **Fault Tolerant Architecture**

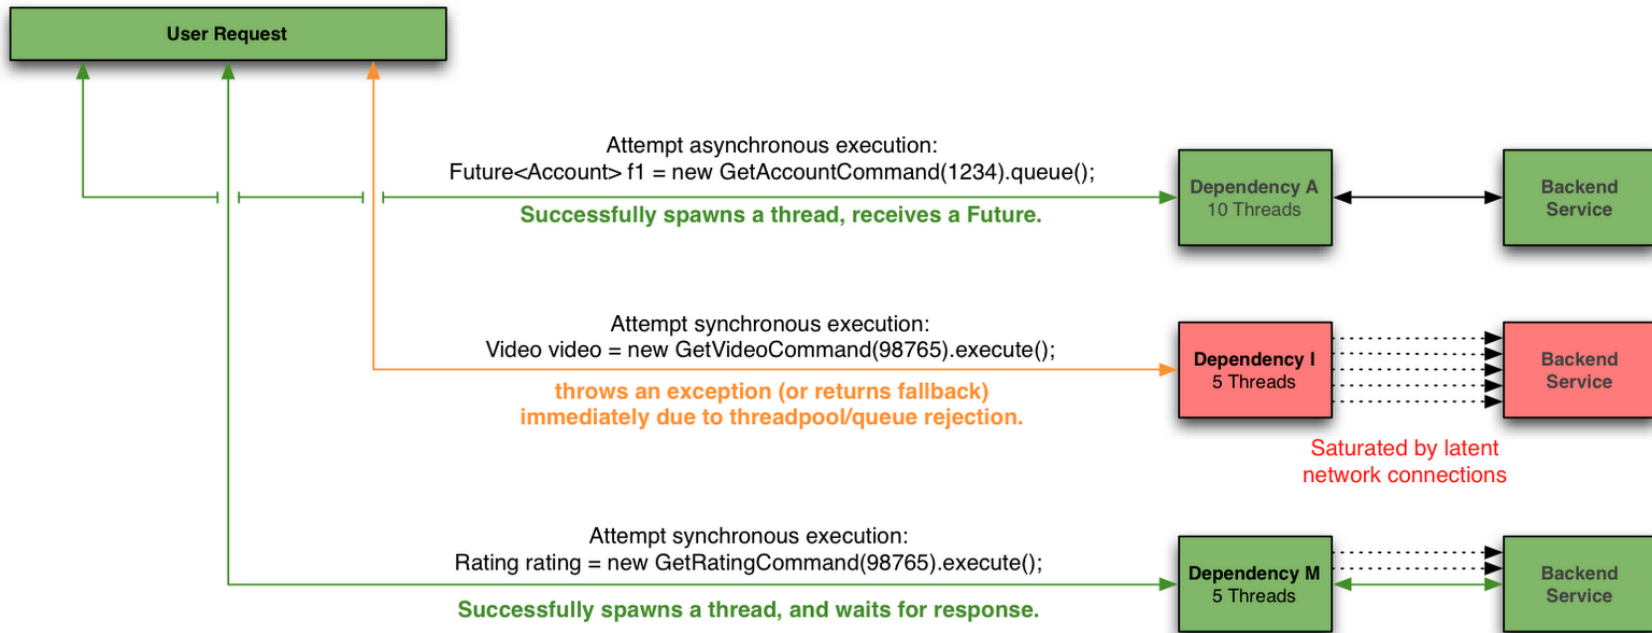
Resiliency/Availability



HYSTRIX
DEFEND YOUR APP

NETFLIX

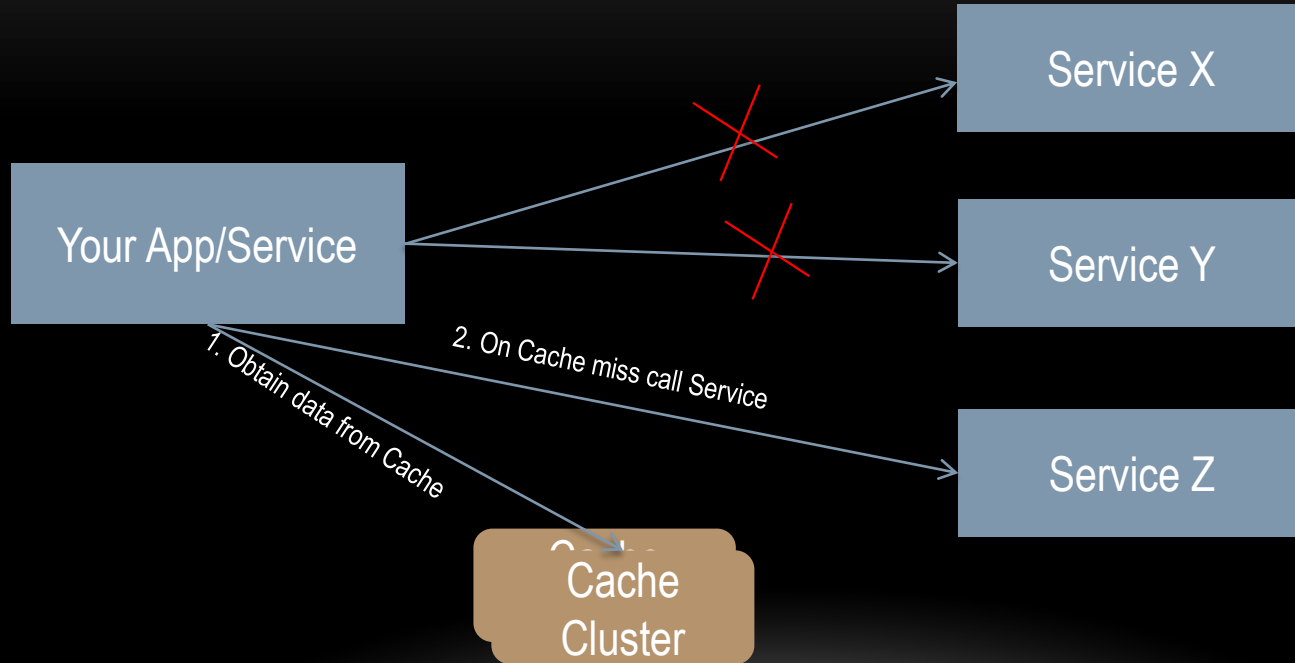
OSS



Circuit Breaker, Retries, Bulk Heading and Fallbacks

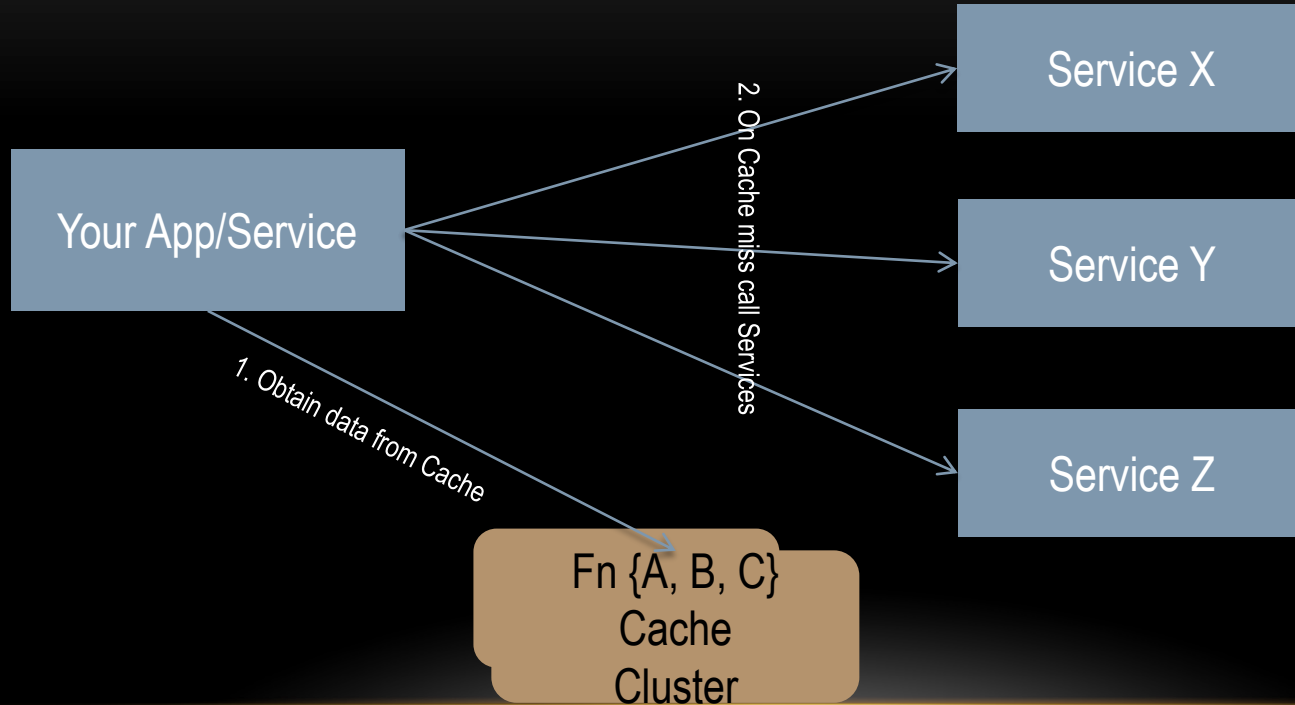
HANDLING FAN OUTS

SERVER CACHING

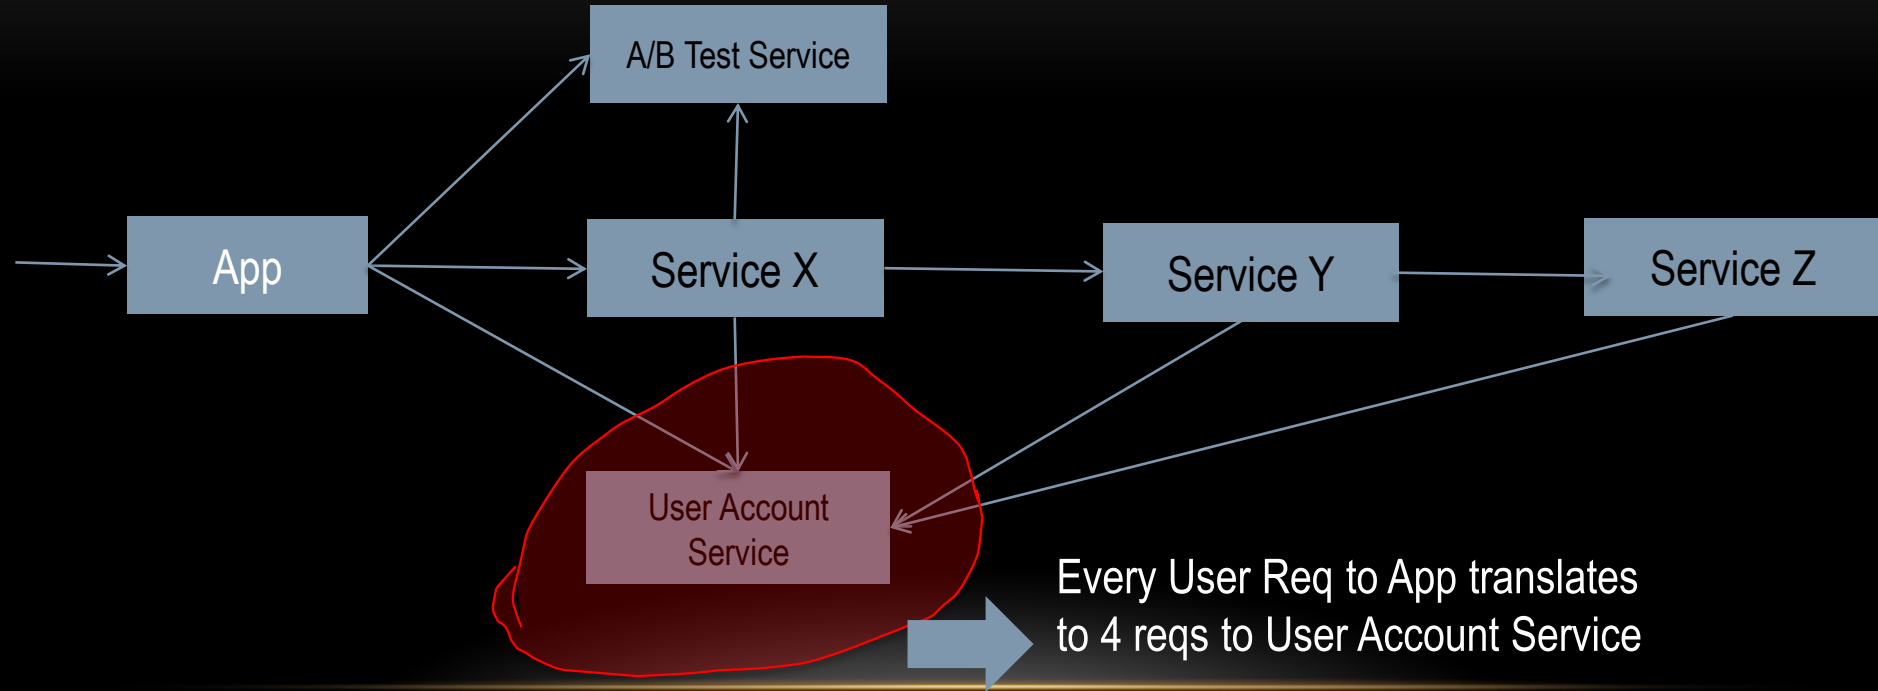


Tip: Config your TTL based on flexibility with data staleness!

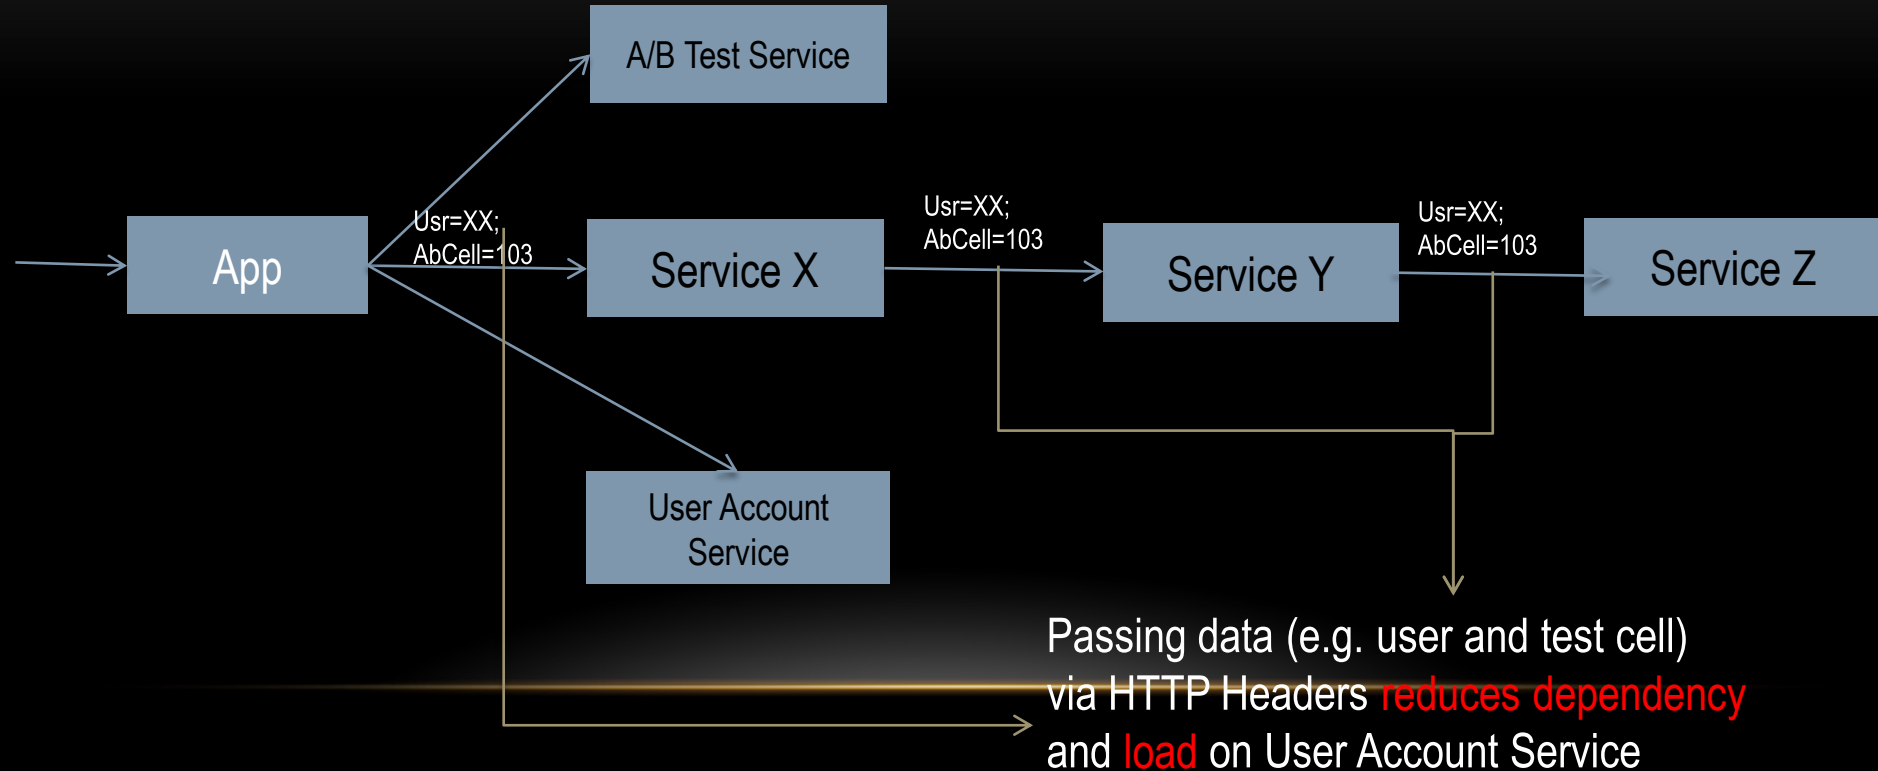
Composite (Materialized View) Caching



BottleNecks/HotSpots



Tip: Pass data via Headers



TEST RESILIENCY (of Overall MicroServices)

- There are only two things certain in life*
 - Death
 - Taxes



* Benjamin Franklin

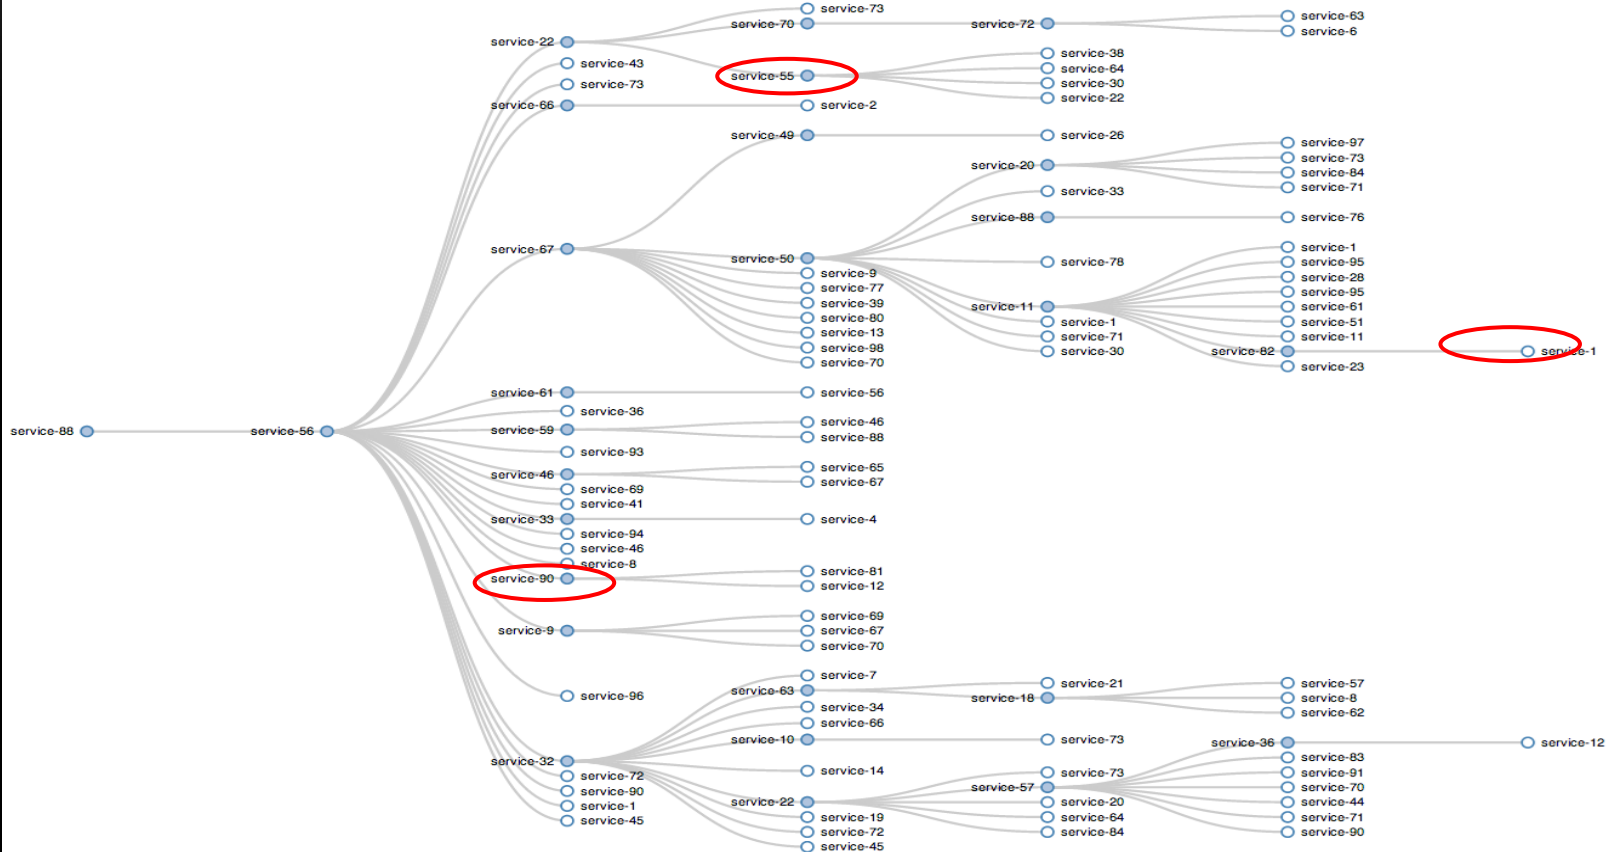
- There are only **three** things certain in life*
 - Death
 - Taxes
 - **Outages in Production**

* Inspired by Benjamin Franklin

Best Practices contd..

- Test Services for Resiliency
 - Latency/Error tests (via Simian Army)
 - Dependency Service Unavailability
 - Network Errors
-

Test Resiliency – to dependencies



TEST RESILIENCY


Use Simian Army <https://github.com/Netflix/SimianArmy>



BEST PRACTICES - SUMMARY


- Isolate your services (Loosely Coupled)
- Use Client Side Smart LoadBalancers
- Dependency Calls
 - Guard your dependency calls
 - Cache your dependency call results
 - Consider Batching your dependency calls
 - Increase throughput via Async/ReactiveX patterns
- Test Services for Resiliency
 - Latency/Error tests (via Simian Army)
 - Dependency Service Unavailability
 - Network Errors

Tools of the Trade



AUTO SCALING

- Use AWS Auto Scaling Groups to automatically scale your microservices
 - RPS or CPU/LoadAverage via **CloudWatch** are typical metrics used to scale



USE CANARY, RED/BLACK PUSHES

- NetflixOSS Asgard helps manage deployments

The screenshot displays the Asgard web interface for managing a cluster named 'obiwan'. The browser address bar shows 'asgardprod/us-east-1/cluster/show/obiwan'. A green box highlights the cluster name in the header, with a green annotation 'This cluster contains two ASGs'. The interface includes a navigation bar with icons for Home, App, AMI, Cluster, ELB, EC2, SDB, SNS, SQS, RDS, and Task. The main content area is titled 'Manage Cluster of Sequential Auto Scaling Groups' and includes a recommendation: 'Recommended next step: Switch traffic to the preferred group, then delete legacy group'. Two ASG groups are shown: 'obiwan-v063' and 'obiwan-v064'. For 'obiwan-v063', a green box highlights the 'OUT_OF_SERVICE' status, with a green annotation 'No traffic on old version'. For 'obiwan-v064', a green box highlights the 'UP' status, with a green annotation 'Live traffic on new version'. A 'Create Next Group' section on the right shows the configuration for 'obiwan-v065', including AMI Image ID, Instance Type, and Instance Counts.

asgardprod/us-east-1/cluster/show/obiwan This cluster contains two ASGs

ASGARD prod us-east-1 CMC

Home App AMI Cluster ELB EC2 SDB SNS SQS RDS Task

Manage Cluster of Sequential Auto Scaling Groups

Recommended next step: Switch traffic to the preferred group, then delete legacy group

obiwan-v063

Launch and Terminate are disabled

Resize to 9 min / 12 max

Delete Disable Enable

9 instances grouped by state

Count State Build ELB Disc

InService 580 OUT_OF_SERVICE

No traffic on old version

obiwan-v064

Resize to 9 min / 12 max

Delete Disable Enable

9 instances grouped by state

Count State Build ELB Disc

InService 583 UP

Live traffic on new version

Create Next Group: Advanced Options

obiwan-v065

AMI Image ID: 179123456789/obiwan-41.2-1417301

Filter

Show more AMIs

Instance Type: m1.large \$230.400/mo

Filter

Instance Counts: Min: 9 Desired: 9 Max: 12

After launch: ☒ Wait for Discovery health check pass

Create Next Group obiwan-v065



Service Dependency Visualization



MicroServices at Netflix



SERVICE DEPENDENCY GRAPH

How many **dependencies** does **my service** have?

What is the **Call Volume** on my Service?

Are any Dependency Services running **Hot**?


What are the **Top N Slowest** “Business Transactions”?

What are the **sample HTTP Requests/Responses** that had a 500 Error Code in the last 30 minutes?

SERVICE DEPENDENCY VISUALIZATION

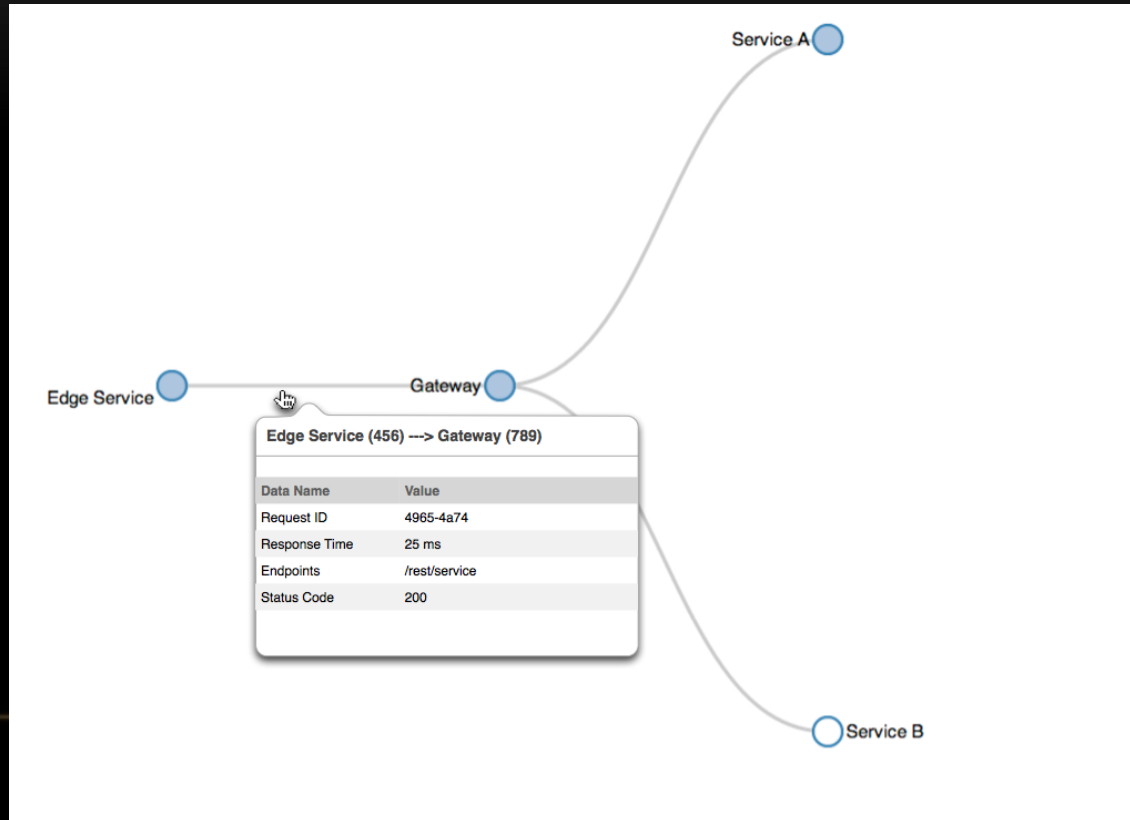


You →

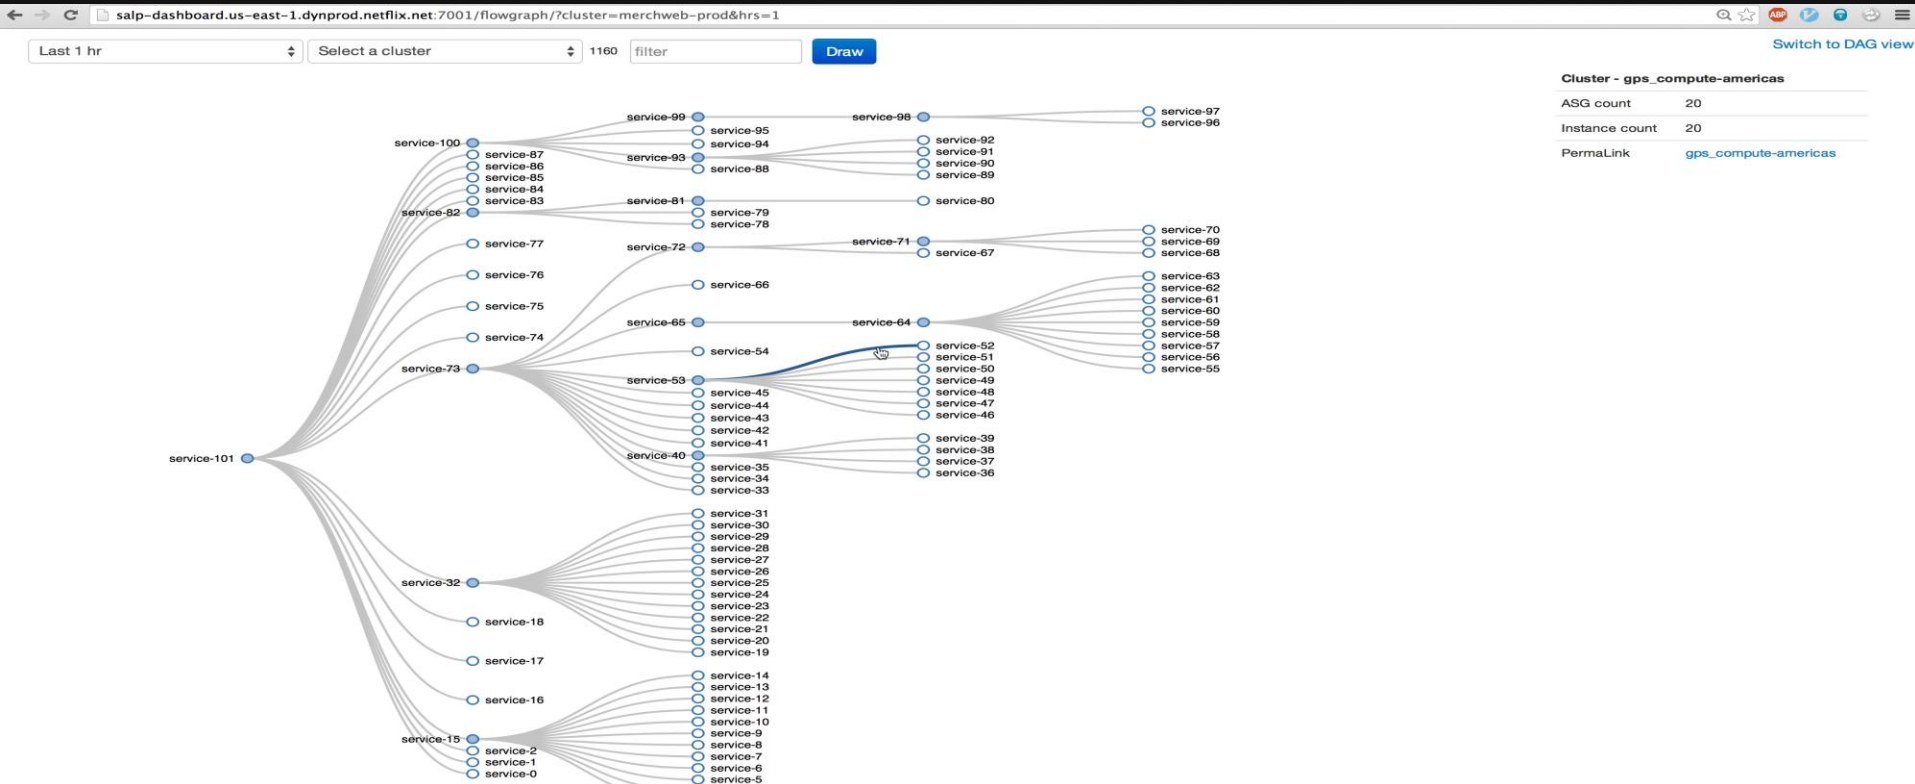


Your Service Dependency Graph

Service Dependency Visualization




Dependency Visualization



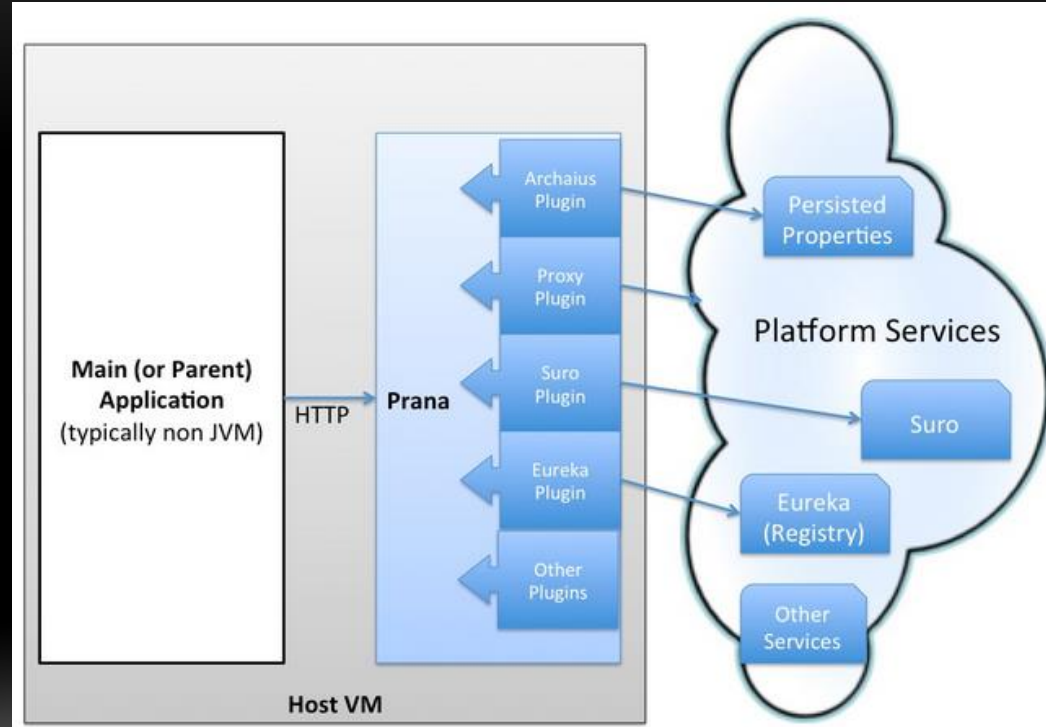
Polyglot Ecosystem

Homogeneity in A Polyglot Ecosystem



TIP: USE A SIDECAR

- Provides a common homogenous Operational/Infrastructural component for all your **non-JVM** based MicroServices



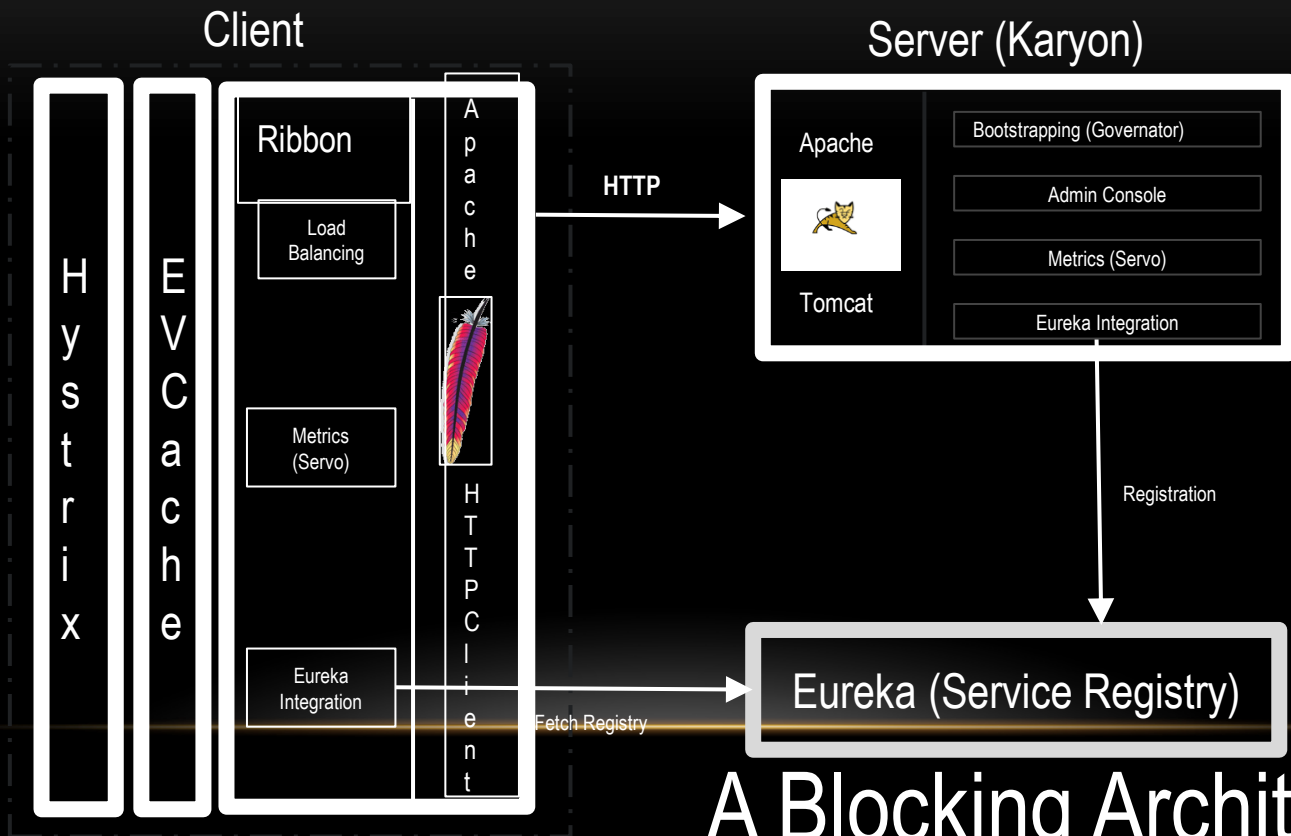
Prana Open Sourced!

- Just this morning!
- <http://github.com/netflix/Prana>

Inter Process Communication

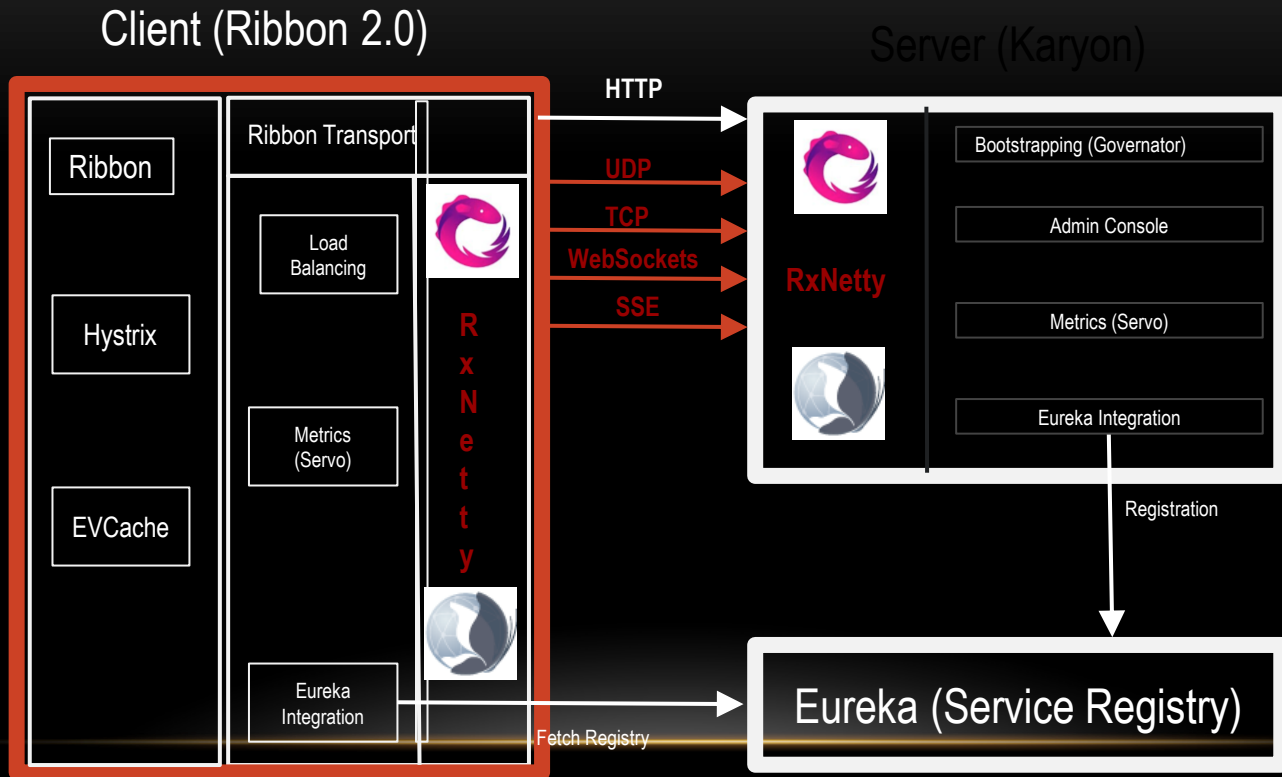


Netflix IPC Stack (1.0)




A Blocking Architecture

Netflix IPC Stack (2.0)



A Completely Reactive Architecture

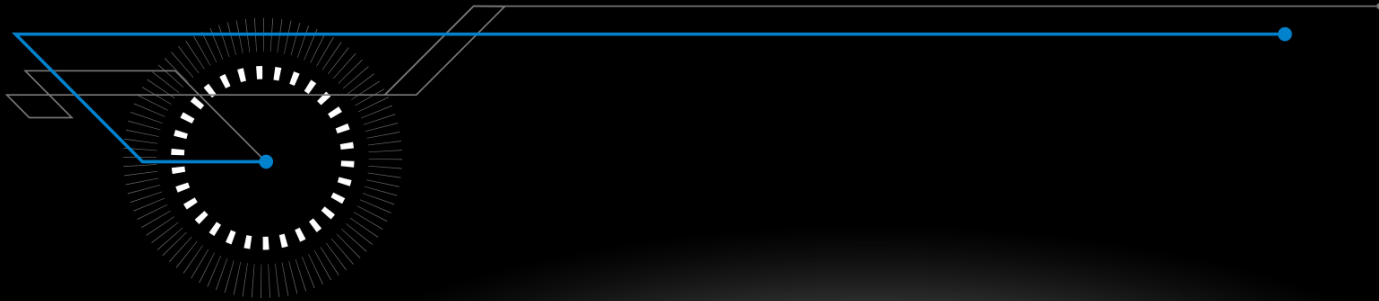
Performance – Throughput



Bounded Thread model (Tomcat) vs Reactive Async (RxNetty)

Details: <http://www.meetup.com/Netflix-Open-Source-Platform/events/184153592/>

NetflixOSS



LEVERAGE NETFLIXOSS

NETFLIX **OSS** Netflix Open Source Software Center

Repositories Powered By NetflixOSS



Getting Started

Welcome to the Netflix Open Source Software Center. To begin, we recommend working with our **RSS Reader application**. See [this walkthrough](#) on Answers For AWS to get up and running quickly.

After you've tackled that, check out the **IBM ACME Air** and **Flux Capacitor** apps.

Also, be sure to join our mailing lists and follow us [@NetflixOSS](#) to stay up to date.

Our Repositories  Thumbnail View  List View 42 public repos 175 members


Availability

HYSTRIX



SIMIANARMY


TURBINE


Cloud Management

ICE


ASGARD


FRIGGA


GLISTEN

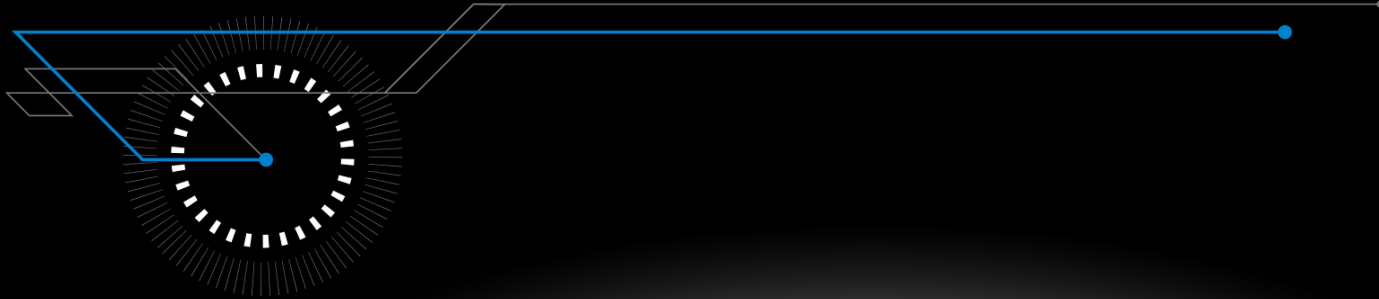

<http://netflix.github.co>

NETFLIX

OSS

- **Eureka** – for Service Registry/Discovery
- **Karyon** – for Server (Reactive or threaded/servlet container based)
- **Ribbon** – for IPC Client
 - And Fault Tolerant Smart LoadBalancer
- **Hystrix** – for Fault Tolerance and Resiliency
- **Archaius** – for distributed/dynamic Properties
- **Servo** – unified Feature rich Metrics/Insight
- **EVCache** – for distributed cache
- **Curator/Exhibitor** – for zookeeper based operations
- ...

Takeaways



Takeaways

- Monolithic apps – good for small organizations
- MicroServices – have its challenges, but the benefits are many
 - Consider adopting when your **organization scales**
 - Leverage Best Practices
 - An Elastic Cloud provides the **ideal** environment (Auto Scaling etc.)
 - NetflixOSS has many libraries/samples to aid you

Questions?